Image Compression with Autoencoder

Replication of work in **Deep Convolutional AutoEncoder-based Lossy Image Compression**

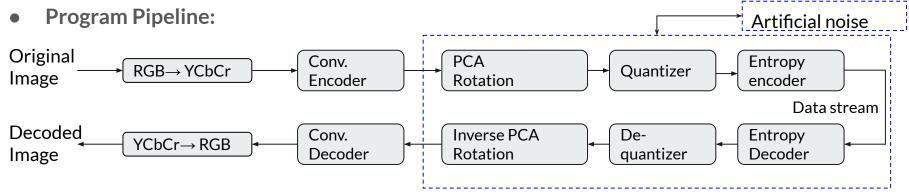
Tian Qiu, Jiawen Zhang ECE, UC San Diego

Background and motivation

- Developing image compression standards: time-consuming
- Deep-learning models can capture complex image features
- Researchers are starting to explore image compression with deep learning
- Works include deep-learning based tools and end-to-end models
 - Toderici, 2017: Recurrent neural network(RNN)
 - o Theis, 2017 : Autoencoder
 - Katto, 2018: Autoencoder with PCA for higher compression rate (Our reference)
- Remaining challenges: Approximating quantization error, computation complexity, limited generalization ability
- Motivation: Try to replicate a work
 - Learn from doing

Overview of reference paper

- Objective: Color image compression with convolutional autoencoder(CAE)
- Main contributions
 - Replace forward/inverse transform with symmetric CAE.
 - After encoder, use PCA to generate more compact information representation



- Training stage: Short-circuit the pipeline, add artificial noise between convolutional encoder and decoder to approximate quantization error
- Application stage: Apply all blocks in series

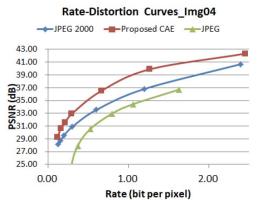
Overview of reference paper

Dataset:

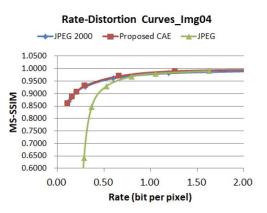
- Training/Val set: A subset of ImageNet (5500 images)
 Original images cropped into 128x128 patches
- Test set: Kodak Lossless Image dataset
 All images in 768x512 size. 24-bit color PNG file.
- Metrics and main results
 - Compression performance: PSNR/SSIM
 - Runtime performance

TABLE I: Average running time comparison.

| Codec | Time (s) |
|--------------------------|----------|
| JPEG | 0.39 |
| JPEG2000 | 0.59 |
| Balle's work[5] with CPU | 7.39 |
| Propose CAE with CPU | 2.29 |
| Propose CAE with GPU | 0.67 |







Test dataset: True Color Kodak Images

Project Overview

- **Objective:** Replicate work in Katto paper
- **Result:** Partially finished the model proposed in the paper
- Main work:
 - Implemented and trained convolutional autoencoder and decoder
 - Implemented PCA analysis to reduce information redundancy
 - Implemented integer quantizer and dequantizer
 - Analysis of model performance on test dataset
- Missing part: Entropy encoding with Adaptive BAC

Implementation details: Setups and dataset

• Programming environment:

UCSD datahub server with GPU. Python 3 and TensorFlow 2.1 environment

• Main packages:

Image processing: opency-python, Pillow

Data processing: numpy, scikit-learn, scipy

Train/Val dataset:

Imagenette: A small subset of ImageNet, consisting of ten easily distinguishable classes.

Total images = 10 classes x 300 pics/class = 3000

Use 80-20 split \rightarrow 2400 training images, 600 test images

Test dataset:

Same as reference(Kodak dataset). 24 images of various topic with rich content.

Implementation details: Image pre-processing

• Image resizing
For train/val sets, resize all images to 128x128.





















Test image partition:



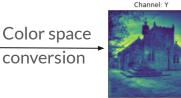
Image partition

For test set, all images can be partitioned to 128x128 blocks.

Encode each block individually

Color space conversion
 RGB→ YCbCr with opency



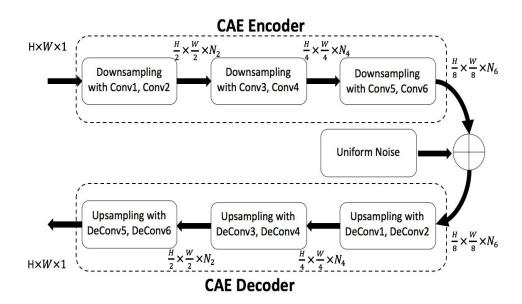


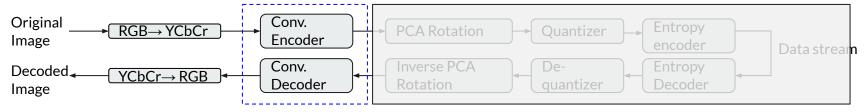




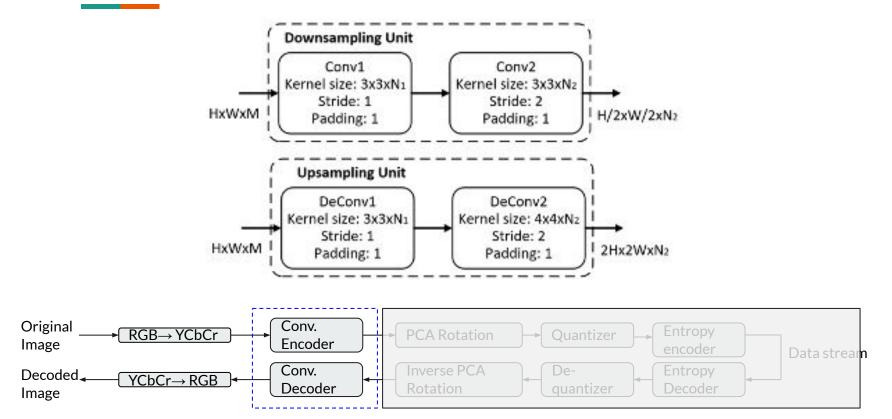
Implementation details: CAE Architecture

- Symmetric CAE network using convolution and deconvolution filters.
- Input size: 128x128x1.
- Number of filters: 32, 32, 64, 64, 64, 32.
- Add uniform noise ranging from -2^-10 to 2^-10 to replace quantization.

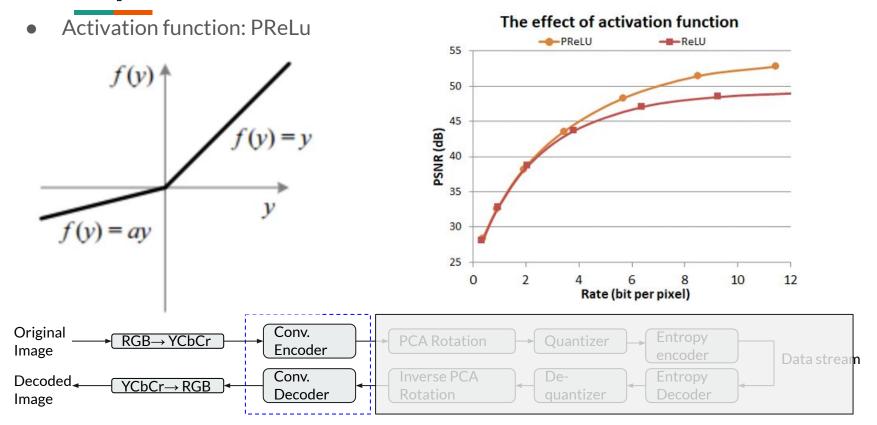




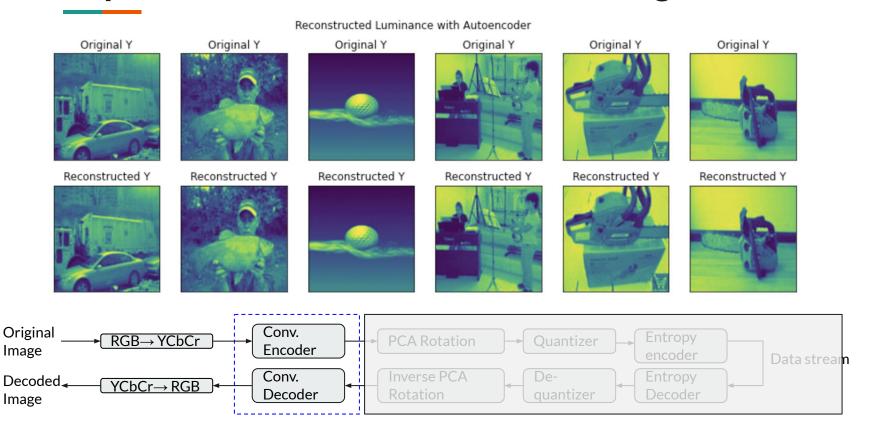
Implementation details: CAE Architecture



Implementation details: CAE Architecture



Implementation details: CAE Training



Implementation details: CAE Training



Reconstructed Image





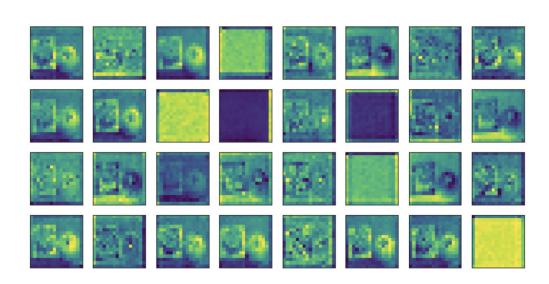
Implementation details: PCA rotation

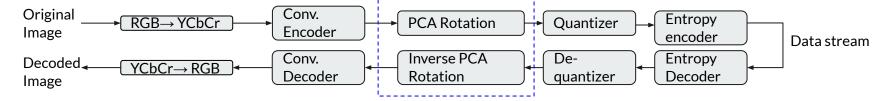
Input:

16x16x32 Encoder output

Output:

16x16x32 data with a more energy-compact representation

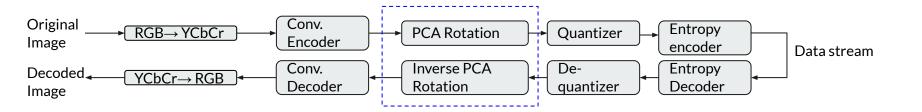




Implementation details: PCA rotation

Algorithm

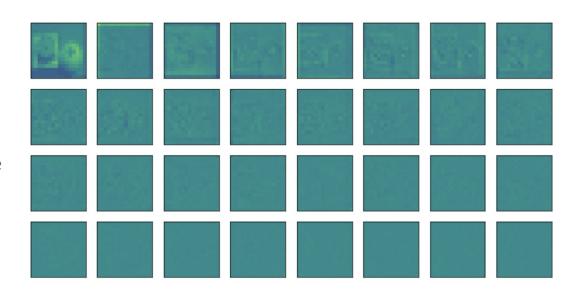
- 1. Reshape the input data from 16x16x32 to 256x32.
- 2. Compute the covariance matrix of the reshaped data.
- 3. Compute the 32 eigenvectors of the covariance matrix.
- 4. Stack the 32 eigenvectors in columns to form matrix *U*.
- 5. Do matrix multiplication
- 6. Reshape the data back to 16x16x32

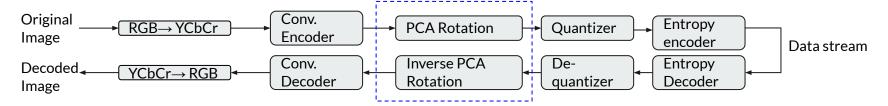


Implementation details: PCA rotation

• Feature maps are sorted in descending order.

The first feature map has the most information.





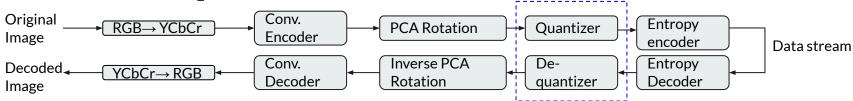
Implementation details: Integer quantizer

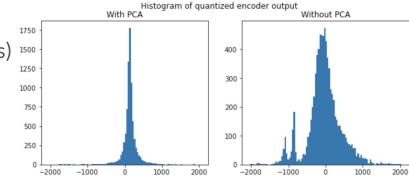
- Input: Encoder output after PCA rotation. Size: [n_samplesx16x16x32] per channel (3 channels)
- Output: Quantized integer representation of input data. Same size
- Parameter: Number of precision = B

 B controls the granularity of uniform quantizer.

 Input data are re-normed and rounded to range [-2^(B-1), 2^(B-1)+1].
 - Example: When B = 12, quantization results ranges from -2048 to 2047
- Formula: $y' = [y \cdot 2^{B-1}]$

$$\tilde{y} = \frac{y'}{2^{B-1}}$$





Implementation details: Entropy encoding

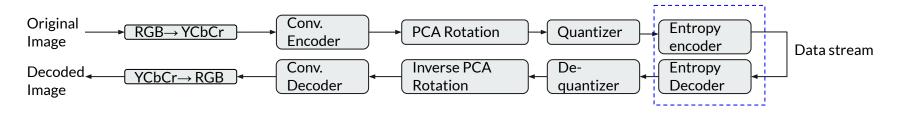
- This part is unfinished.
- Approach in reference paper: Compress quantized integer with Adaptive Binary Arithmetic Coder(CABAC)
- Our approach:

Do not actually compress. Instead, estimate required bitrate with source entropy.

H = discrete entropy of quantized integer sequence for one channel (size: <math>16x16x32)

Rough estimate of bit per pixel(bpp):

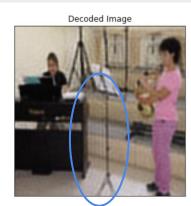
$$\text{bpp} \geq \frac{\text{estimated total encoded bits}}{\text{number of pixels}} = \frac{16 \times 16 \times 32 \times (H_Y + H_{Cr} + H_{Cb})}{128 \times 128}$$



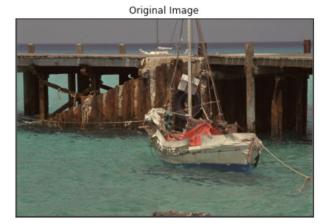
Experiment results: Visual

- Validation images:
 Blurry in general
 Obvious ringing/'staircase' artifact
- Test images:
 Less obvious artifact
 Random spots

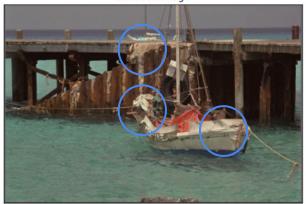




Precision = 12, PSNR = 29.1, SSIM = 0.91







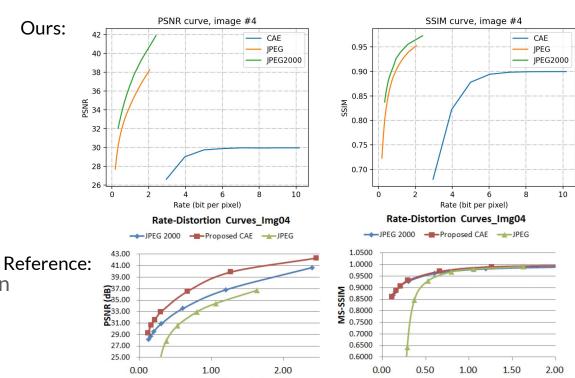
Precision = 12, PSNR = 29.6, SSIM = 0.87

Experiment results: Quality Metrics

Rate-distortion curve:
 PSNR limited under 30
 SSIM limited under 0.9
 Bpp is high

Possible reasons

- CAE needs better generalization ability to work on large test set.
- PCA and quantization can be more compact



JPEG implemented with Python Pillow package: https://pillow.readthedocs.io/en/stable/handbook/image-file-formats.html#jpeg
JPEG2000 implemented with Python Glymur package on top of OpenJPEG Library: https://glymur.readthedocs.io/en/latest/

Limitations, future work, conclusion

• Limitations:

Not fully replicating the paper

• Future work:

Include entropy encoding/decoding

Use larger dataset

Explore more compact structure of autoencoder

Conclusion:

Implemented a image compression system with convolutional autoencoder.

Demonstrated the potential of CAE to compress image.

Demonstrated the ability of PCA to generate compact information representation.

Notes on project code

- The whole project is in a jupyter notebook
- We used an isolated conda environment for the project
 - A Markdown file create_285_env.md is included in the submission for reference.
 - A list of required packages is in requirements.txt
 - Use the commands in the Markdown file to recreate the project environment
- We will quickly go through the code after the presentation

Thank you!

• References:

Z. Cheng, H. Sun, M. Takeuchi and J. Katto, "Deep Convolutional AutoEncoder-based Lossy Image Compression," 2018 Picture Coding Symposium (PCS), San Francisco, CA, 2018, pp. 253-257, doi: 10.1109/PCS.2018.8456308.

Toderici, George, et al. "Full resolution image compression with recurrent neural networks." *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*. 2017.

Theis, Lucas, et al. "Lossy image compression with compressive autoencoders." ICLR. 2017

Uniform noise layer: https://github.com/RitwikGupta/Keras-UniformNoise

Test dataset: True Color Kodak Images

Imagenette: A small subset of ImageNet developed by fast.ai: https://github.com/fastai/imagenette

Python Pillow package: https://pillow.readthedocs.io/en/stable/handbook/image-file-formats.html#jpeg

Python Glymur package: https://glymur.readthedocs.io/en/latest/